

Audit research notes

Matteo Franzil, Valentino Armani, Luis Augusto Dias Knob, Domenico Siracusa

July 10, 2024

Contents

1	AIs	1
1.1	DONE SoA research on log clustering, graph representation, and mapping	1
1.2	DONE Implement a rough parser and filter out background noise	1
1.3	DONE Investigate difference between <code>kubect1</code> and other tools	2
1.4	SCHE Data generation and collection	2
1.5	SCHE Implementation of an ML model for learning behaviours	2
1.6	SCHE Implement the final visualizer	4
1.7	SCHE Strengthen the motivation for the project	4
1.8	SCHE Retrieve additional datasets	4
1.9	Other things	4
2	Strengthening the motivation for the project	4
3	Behaviour analysis of the audit logs	5
3.1	Generic	5
3.2	Namespaces	5
3.3	Other quirks	5
3.4	Impersonation	6
3.5	Control plane components	6
4	Future work	6

1 AIs

1.1 **DONE** SoA research on log clustering, graph representation, and mapping

1.1.1 **DONE** Inspect SoA for what other people did in the field

- We did some SoA research to understand if other works performed log clustering and pattern matching. The results are quite unsatisfactory: most works - obviously - focus on NLP processing for the log content, few attempt to actually clusterize multiple lines.

[1] does so in a Hadoop context, a distributed MapReduce system where failures are common. Starting from the assumption that events are limited in quantity, they vectorize sequences over N dimensions (N different events) and clusterize them using similarity.

[2] LogMiner (first author also wrote [1]) also does something similar to what we want to do: transform log sequences into *behaviours*, depending on the context the various items acted upon. They again use inverse document frequency to assign weight to logs in each cluster, and partitions behaviours using PIDs. The amount of behaviours is, again, quite limited.

1. Q. Lin, H. Zhang, J.-G. Lou, Y. Zhang, and X. Chen, “Log clustering based problem identification for online service systems,” in Proceedings of the 38th International Conference on Software Engineering Companion, in ICSE ’16. New York, NY, USA: Association for Computing Machinery, May 2016, pp. 102–111. doi: 10.1145/2889160.2889232. Available: <https://dl.acm.org/doi/10.1145/2889160.2889232>. [Accessed: Jun. 18, 2024]
2. H. Zhang, L. Cai, L. Zhao, A. Yu, J. Ma, and D. Meng, “LogMiner: A System Audit Log Reduction Strategy Based on Behavior Pattern Mining,” in MILCOM 2022 - 2022 IEEE Military Communications Conference (MILCOM), Rockville, MD, USA: IEEE, Nov. 2022, pp. 292–297. doi: 10.1109/MILCOM55135.2022.10017626. Available: <https://ieeexplore.ieee.org/document/10017626/>. [Accessed: May 09, 2024]

1.2 DONE Implement a rough parser and filter out background noise

1.2.1 DONE Finish clustering non-relevant events and implement them in the parser

We now have a parser that filters out background noise and separates control plane traffic from user traffic. The control plane traffic will be unlabeled, while the user traffic will be labeled depending on the behavior.

1.3 DONE Investigate difference between kubectl and other tools

- `kubectl` performs a series of sanity checks before actually sending the request to the API server. This is done by the client itself, for example, getting the namespace and ResourceQuotas before creating an object.
- On the other hand, other behaviours are managed by the KCM, for example, creating a secret for a ServiceAccount, or creating a ConfigMap for the root CA.

1.3.1 DONE Try to use curl or an external library

- Does not really change much, user agent apart. Some behaviours are still triggered by the KCM, and API calls made have no effect on his decisions.

1.3.2 FAIL Inspect code to see how api calls are made

- Utter chaos, depending on the command the API calls are different: sometimes it uses the API objects, sometimes raw HTTP requests... plus it uses command builders that are not easy to follow.

1.4 SCHE Data generation and collection

1.4.1 DONE Take notes of weird behaviours

Have a look at behaviour analysis for a detailed list of weird behaviours that we have found in the logs.

1.4.2 SCHE Collect datasets

- **DONE** Category 1
- **DONE** Category 2
- **TODO** Category 3

1.5 SCHE Implementation of an ML model for learning behaviours

1.5.1 DONE Label definition

To ease in the labeling and recognition of the data, we have designed a dynamic labeling system that allows us to label data in a more efficient way.

Labels are 22-bit-long integers that encode the following information:

- The type of object that is being acted upon, plus any subtypes
- The action that is being performed on the object
- Whether the object is namespaced or not and if the action is performed on a list of objects or a single object

Labels are not per-line, but rather per-behaviour. For example, all the things that happen when a namespace is created are labeled with the same label: this very label encodes the fact that a namespace (type of object) is being created (action) and that the action is performed on a single object (single object).

To simplify, we have defined a set of "simple" actions that are used to label the data: **create**, **delete**, **get**, **watch**, **update**. These actions combine verbs that are present in the audit logs such as **delete** and **deletecollection** into a single action, allowing us to abstract away from the actual verb used in the log and reduce the label space.

To further reduce the label space, the API specification was used to identify invalid combinations of actions and objects, reducing the number of possible labels to 490.

This label system allows us to easily identify patterns in the data in a quick and efficient way: the parser will be able to inspect the logs and immediately understand what is happening without having to look at the actual log content.

1.5.2 DONE Creation of a pipeline

- The result of our initial work is a pipeline that, taking

raw audit logs as an input, does the following:

- Filters out background noise (e.g. `/apis`, `/healthz`, `/version`)
- Identifies simple control plane behaviours (e.g. service account

token renewal), tags it with `cplabel` to `true` and assigns a label.

- IMPORTANT! This is a stark change to last time. We decided to include all control plane API traffic in the datasets.
- Prompts the user to start the manual labeling process on the filtered data.
- Re-merges the labeled data with the original data.

The data is then ready to be used for training a model.

1.5.3 DONE Model choice

- **FAIL** KNN
 - To avoid diving on difficult and complicated models I first tried something simple (and stupid): K Nearest Neighbors. KNN is a quick and dirty way to classify data, but it was almost instantly clear that it was not going to work.

- **FAIL CNN**

- I then tried with CNNs, with better yet imperfect results.
 - * Accuracy was somewhat OK (90/94%), but the model size quickly spiraled out of control and so were the training times.
 - * Convolutional layers usually are interleaved with pooling layers to reduce the size of the data, but in our case the pooling layers made more harm than good. Trying to remove them made the model even bigger, keeping them jeopardized the accuracy.
 - * The model was able to pick up on some patterns, but not to generalize well: overfitting was almost guaranteed.

- **DONE LSTM**

- I then decided to go back to study RNNs and I shortlisted two promising candidates: LSTM and variational autoencoders.
- So far, LSTM performance has been extremely promising: over 40 randomized attempts, the model has been able to reach an F1 score of 0.98 and accuracy, recall, and precision all over 0.97.
 - * P.S. all these metrics have been calculated with the worst possible odds in our favour, e.g., precision and recall with macro averaging.

- **SCHE Improve the model**

- Work has been done on adapting the model to the new Cat2 data
- Exploratory research on WINDOW_{SIZE} (40 -> 60), training splits, and hyperparameters.
- The model also demonstrated to work with completely randomized sequences with Cat1. We need to study this deeper.
- Further work is needed with refining features and encodings.

1.6 **SCHE Implement the final visualizer**

- Valentino is working on this.

1.7 **SCHE Strengthen the motivation for the project**

See motivation for a detailed explanation of the motivation behind the project. This section is WIP.

1.8 **SCHE Retrieve additional datasets**

1.8.1 **DELE Ask PoliTO for a better dataset**

1.8.2 **TODO Generate random data by hand**

- The data we collected in our dataset generation was joined by a "dataset" me and Luis generated by randomly and repeatedly executing commands in a Kubernetes cluster. This dataset must still be labeled

1.9 **Other things**

1.9.1 **FAIL Scrapped: learn templates from the logs**

- Idea: use a ML model to generalize and attempt to match 'templates' of behaviour in the logs. For example: we already have an idea of how namespaces are created, deleted, and so on. We could use a model to learn these templates and then match them in the logs. If no match is found, then an attempt could be made to generate a new template based on the most similar ones.

2 Strengthening the motivation for the project

- Signature-based detection is not enough. To demonstrate so, we can use the example of Falco. Glossing over the nightmare that is installing it, Falco comes with a pre-set of rules that are fairly basic. For example, rules are triggered when SAs are created and deleted, secrets are accessed, and so on. However, these rules are exceedingly verbose and alerts are triggered for use cases that are not malicious at all. For example, Flannel renewing SA tokens, or the kube-controller-manager creating a secret for a SA when a namespace is created, are all examples of false positives that are triggered by Falco.
- ValidatingWebhookConfiguration is a powerful tool that can be used to validate requests to the API server. However, it is not used in the default installation of Kubernetes, and it is not easy to set up. Furthermore, it is not easy to understand what is happening when a request is denied, as the logs are not very verbose.

3 Behaviour analysis of the audit logs

3.1 Generic

- **watch** requests are long-running and are usually made by control plane components. Since the **ResponseStarted** and **ResponseComplete** events can be in different times, we chose to **remove** the **ResponseStarted** event from the logs.
 - Indeed, for us it is important to have the **requestReceivedTimestamp** that is present in all events anyway. With that, we can understand when the request was made and when it was completed.
 - **ResponseStarted** appears in **watch** requests, in **kubectl exec** requests, and in **delete** requests. In the latter case, **ResponseComplete** is sent once the object is actually deleted.
- **kubectl** requests are internally translated into a set of API calls that are made by the **kube-apiserver**. Further subsections detail such calls and how they can help us derive templates for the parser.

3.2 Namespaces

- Every time a namespace is created, the following things happen:
 - A **list** request is made to **/api/v1/namespaces/{}/resourcequotas** by the **kube-apiserver** for no apparent reason
 - * To add insult to injury, usually this line is out of order in the log
 - The **kube-controller-manager** gets the **service-account-controller** SA and generates a token for it
 - The SA **root-ca-publisher** puts the **root-ca.crt** in a ConfigMap in the namespace
 - The **service-account-controller** creates a SA called **default** in the namespace
- When a namespace is deleted, all subresources are deleted with a **deletecollection** request followed by a **list**. Exceptions include **Pods**, that weirdly are deleted by a **list**, then **deletecollection**, then **list** again. Also, when **configmaps** are deleted, the **kube-controller-manager** (and not the **root-ca-cert-publisher**) tries to create again the CM that was just deleted, but is rejected because the namespace is being terminated.

3.3 Other quirks

- When a secret is created for a SA the secret is instantly patched by the kube-controller-manager, which is in charge of actually generating that secret along with an object like

```
{"data":{"ca.crt":"...","namespace":"...","token":"..."}
```

where the three fields are base64 encoded.

- Secrets can be traced back to the Pod they have been mounted into by either looking at the spec or monitoring for `watch` requests by the node that is hosting the pod.
- Users with access to `kubectl exec` can access the secrets of the pods they are exec-ing into using `printenv`.
- ResourceQuotas can be used for more than just limiting resources, they can also be used to limit the number of objects that can be created. This also applies to things like `serviceaccounts`: thus, it seems that every time a potentially limitable object is created, a `get` request is made to `/api/v1/{object}` to assess the current state.
- When CSRs are being approved or denied, the `update` verb is used and the actual decision is in the payload of the request. This makes it less straightforward to understand if the request was allowed or denied.
- When old-style ServiceAccount tokens are used, instead of the new CSR style Users, the API server detects the *offending* action and adds a label `kubernetes.io/legacy-token-last-used` to the Secret.

3.4 Impersonation

- Users with access to the `impersonate` verb and one of `users`, `groups`, and `serviceaccounts` can impersonate other identities.
- This is correctly reflected in the audit logs, but could be a problem implementation-wise.

3.5 Control plane components

- Control plane components' "background noise" has been filtered out and the details are available in [queries.org](https://kubernetes.io/queries.org).

4 Future work

- Using KWOK
- Ask for a master's student to create a simulator with KWOK, defining scenarios and gathering the audit log